



NTU COMPUTING CHALLENGE: TEAM OMEGA

Anderson Serangoon Junior College



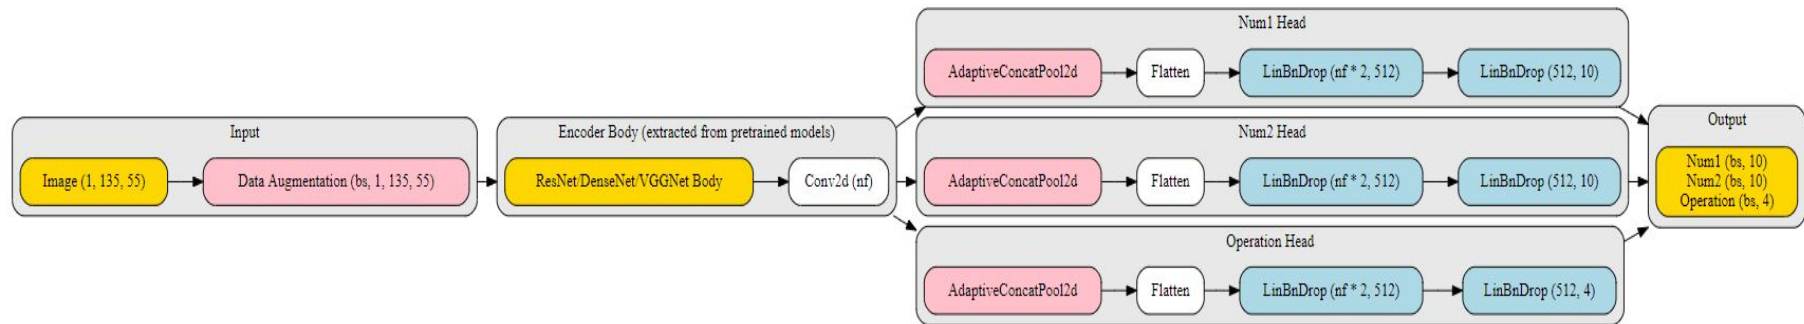
RECAP OF PROBLEM STATEMENT (HANDWRITTEN MATH EXPRESSIONS)


2+7



9.00

MODEL ARCHITECTURE (MULTI-HEAD CNN)





“FASTAI” FRAMEWORK

Howard, J., & Gugger, S. (2020). Fastai: A layered API for deep learning. *Information*, 11(2), 108.



WHAT IS “FASTAI”

fastai is a deep learning library which provides practitioners with high-level components that can quickly and easily provide state-of-the-art results in standard deep learning domains, and provides researchers with low-level components that can be mixed and matched to build new approaches.

The logo for fast.ai, featuring the text "fast.ai" in a white, lowercase, serif font. The text is underlined with a thick white horizontal line.





Making neural nets
uncool again



DATA AUGMENTATION

Data augmentation in data analysis are techniques used to increase the amount of data by adding slightly modified copies of already existing data.






RESIDUAL NEURAL NETWORK (RESNET)

Kaiming He, Xiangyu Zhang, Shaoqing Ren
and Jian Sun (2015). Deep Residual Learning
for Image Recognition



ADDING MORE LAYERS ON TOP OF A NETWORK WILL DEGRADE ITS PERFORMANCE




Source: Original paper



RESIDUAL BLOCK

SKIP CONNECTIONS

The skip connection skips training from a few layers and connects directly to the output.




ADVANTAGE

If any layer hurts the performance of architecture then it will be skipped by regularization.



vanishing/exploding gradient problem



DENSE CONVOLUTIONAL NETWORK (DENSENET)

Gao Huang, Zhuang Liu, Laurens van der Maaten,
Kilian Q. Weinberger (2016), Densely Connected
Convolutional Networks



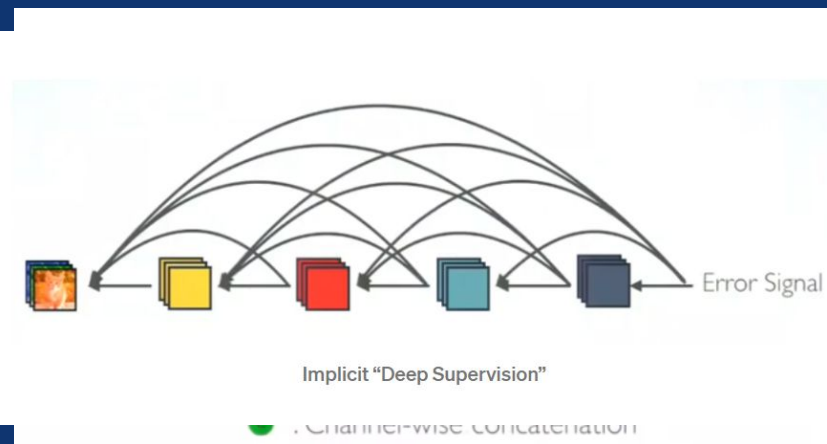
DENSENET

CONCATENATION

Each layer is receiving a “collective knowledge” from all preceding layers.

GROWTH RATE

Network can be thinner and more compact. i.e number of channels can be fewer



PARAMETER AND COMPUTER EFFICIENCY

DenseNet has a smaller size than ResNet

STRONG GRADIENT FLOW

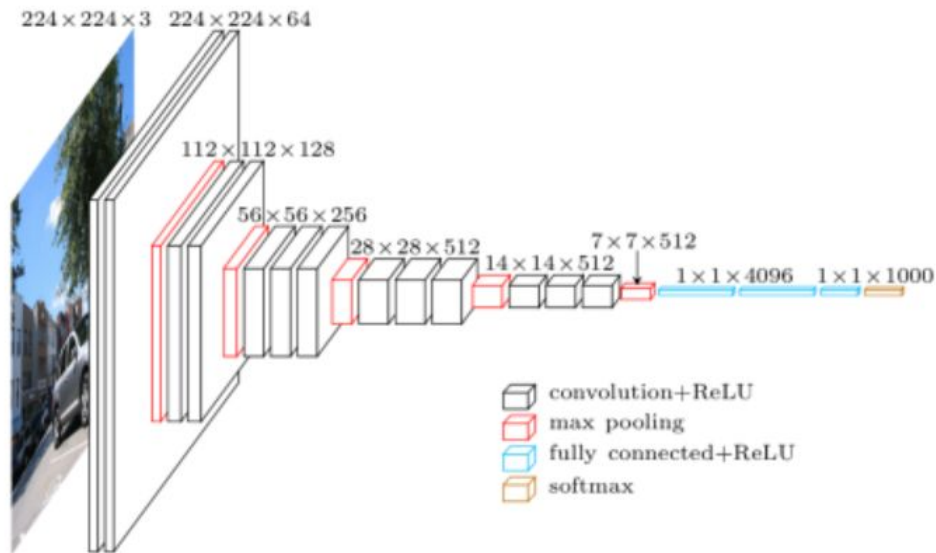
Error signal can be easily propagated to earlier layers more directly.

VISUAL GEOMETRY GROUP (VGG) NET

Karen Simonyan, Andrew Zisserman
(2015), Very Deep Convolutional Networks
for Large-Scale Image Recognition




VGGNET



- uses only 3x3 convolutional layers stacked on top of each other in increasing depth
- reducing volume size is handled by max pooling
- two fully-connected layers, each with 4,096 nodes are then followed by a softmax classifier



VGGNET



By using 2 layers of 3×3 filters, it actually have already covered 5×5 area. By using 3 layers of 3×3 filters, it actually have already covered 7×7 effective area.



Number of parameter are fewer.



Faster convergence and reduce the overfitting problem.






TRANSFER LEARNING

Kafeng Wang, Xitong Gao, Yiren Zhao,
Xingjian Li, Dejing Dou, Cheng-Zhong Xu
(2020), Pay Attention to Features, Transfer
Learn Faster CNN





COMBINING ALL THREE

Encoder Body	Max Accuracy (Validation)
res50	0.9959
dense169	0.9956
vgg19	0.9963
all 3 combined	0.9972

WE GET A HIGH ACCURACY




ONE CYCLE POLICY FOR LEARNING RATE

Leslie N. Smith (2018), A Disciplined Approach to
Neural Network Hyper-Parameters: Part 1 --
Learning Rate, Batch Size, Momentum, and Weight
Decay




EFFECTS OF LEARNING RATES ON TRAINING

Too low




A small learning rate requires many updates before reaching the minimum point

Just right




The optimal learning rate swiftly reaches the minimum point

Too high



Too large of a learning rate causes drastic updates which lead to divergent behaviors

CYCLIC LEARNING RATES



1 -1

Cycle and
Stepsize
Saddle Point




CROSS ENTROPY LOSS FUNCTION

Zhilu Zhang, Mert R. Sabuneu (2018),
Generalized Cross Entropy Loss for Training
Deep Neural Networks with Noisy Labels



CROSS ENTROPY LOSS FUNCTION



ADAM OPTIMIZER

Diederik P. Kingma, Jimmy Ba (2017), Adam: A Method for Stochastic Optimization



HOW DOES ADAM WORK?

1) **ADAPTIVE GRADIENT ALGORITHM (ADAGRAD)**

that maintains a per-parameter learning rate that improves performance on problems with sparse gradients (eg. natural language and computer vision problems)

HOW DOES ADAM WORK?

2) ROOT MEAN SQUARE PROPAGATION (RMSPROP)

that also maintains per-parameter learning rates that are adapted based on the average of recent magnitudes of the gradients for the weight (eg. how quickly it is changing). This means the algorithm does well online and non-stationary problems.



BENEFITS OF USING ADAM

STRAIGHTFORWARD TO IMPLEMENT





BENEFITS OF USING ADAM

COMPUTATIONALLY EFFICIENT





BENEFITS OF USING ADAM

LITTLE MEMORY REQUIREMENTS





BENEFITS OF USING ADAM

**INVARIANT TO DIAGONAL
RESCALE OF THE GRADIENTS**





BENEFITS OF USING ADAM

**WELL SUITED FOR PROBLEMS THAT
ARE LARGE IN TERMS OF DATA OR
PARAMETERS**





BENEFITS OF USING ADAM

**HYPER-PARAMETERS HAVE INTUITIVE
INTERPRETATION AND TYPICALLY
REQUIRE LITTLE TUNING**



OUR BUSINESS IDEA



A microscopic image of a virus particle, likely SARS-CoV-2, showing its characteristic spherical shape with a textured surface and prominent red, spike-like projections. A solid blue rectangular box is centered over the image, containing the text "COVID-19 PANDEMIC" in white, bold, sans-serif capital letters.

COVID-19 PANDEMIC

TIME FOR NEW BEGINNINGS. GET 2 MONTHS FOR ONLY \$8

SUBSCRIBE NOW

TIME

DEMOCRACY UNDER ATTACK

BIDEN'S VICTORY CERTIFIED

WILL TRUMP FACE CONSEQUENCES?

NEWSLETTER

SUBSCRIBE



U.S. • COVID-19

'Where Are the Kids?' School Is Back in Session, but Many Kindergarteners Are Missing



CNA Insider

CNA Lifestyle

CNA938

Singapore

Asia

World

Commentary

Business

Sport

Climate Change

News Clips

V

All Sections



By Chew Hui Min

03 Apr 2020 05:18PM
(Updated: 03 Apr 2020 09:38PM)



Bookmark



Singapore

'Right time' to close schools now, says Education Minister Ong Ye Kung



MATH DOODLE

